

```
In [1]: using DiffEqFlux, DifferentialEquations, DiffEqCallbacks, SciMLBase, Distributions, LuxCore
using Statistics, LinearAlgebra, Plots
using Flux.Data: DataLoader
using Lux, Optimization, OptimizationOptimJL, Optimisers, OptimizationOptimisers, Random, Plots
using ComponentArrays

using OrdinaryDiffEq, Flux, MLDataUtils, NNlib
using Flux: logitcrossentropy
using MLDatasets: MNIST
```

We implement Augmented Neural ODEs to perform classification on MNIST handwritten digit images. The original Neural ODE image classification code (without any augmentation) is adapted from <https://docs.juliahub.com/DiffEqFlux/BdQ4g/1.10.1/examples/Supervised-NN-ODE-MNIST/>. We implement several potential performance improvements such as zero-padding augmentation, input-layer augmentation, temporal regularization, and second-order Neural ODEs. We measure and visualize performance using test accuracy and loss over iterations.

Loading MNIST dataset

```
In [2]: function loadmnist(batchsize = bs)
# function to convert Labels to one-hot format for training
onehot(labels_raw) = convertLabel(LabelEnc.OneOfK, labels_raw, LabelEnc.NativeLabels(collect(0:9)))
# Load MNIST train and test sets
imgs = MNIST.traintensor(1:1000)
labels_raw = MNIST.trainlabels(1:1000)
test_imgs = MNIST.testtensor(1:1000)
test_labels_raw = MNIST.testlabels(1:1000)
# Process images into (H,W,C,BS) batches
x_train = Float32.(reshape(imgs,size(imgs,1),size(imgs,2),1,size(imgs,3)))
x_train = batchview(x_train,batchsize)
x_test = Float32.(reshape(test_imgs,size(imgs,1),size(imgs,2),1,size(imgs,3)))
x_test = batchview(x_test,batchsize)
# Onehot and batch the Labels
y_train = onehot(labels_raw)
y_train = batchview(y_train,batchsize)
y_test = onehot(test_labels_raw)
y_test = batchview(y_test,batchsize)
return x_train, y_train, x_test, y_test
end

bs = 100 # batch size
x_train, y_train, x_test, y_test = loadmnist(bs)
```

Helper functions for training Neural ODE for classification

```
In [9]: # helper function to convert ODE solver output to array that can be fed into next layer of neural network
function DiffEqArray_to_Array(x)
new_array = reshape(x[1], size(x)[1:2])
return new_array
end

# helper function to convert Logit predictions to digit classes (0-9)
classify(x) = argmax.(eachcol(x))

# compute classification accuracy fraction of model on data
function accuracy(model,data; n_batches=10)
total_correct = 0
total = 0
for (x,y) in collect(data)[1:n_batches]
target_class = classify(Flux.cpu(y))
predicted_class = classify(Flux.cpu(model(x)))
total_correct += sum(target_class .== predicted_class)
total += length(target_class)
end
return total_correct/total
end

# compute Logit cross entropy loss of model on data
function total_loss(model, data; n_batches=10)
total = 0
for (x,y) in collect(data)[1:n_batches]
total += logitcrossentropy(model(x),y)
end
return total
end

# ADAM optimizer (will be used for all ANODE variants)
global opt = Flux.Optimise.ADAM(0.05)
```

```
Out[9]: Flux.Optimise.Adam(0.05, (0.9, 0.999), 1.0e-8, IdDict{Any, Any}())
```

Neural ODE variant: temporal regularization

We implement a Neural ODE variant with temporal regularization. The code is adapted from that for the NeuralODE from DiffEqFlux.jl at https://github.com/SciML/DiffEqFlux.jl/blob/1e4614c7c0b0115233f8da09cf23ee37202b830d/src/neural_ode.jl. Our main modification is that there is an extra end_time_regularization parameter, and the endpoint t2 of the ODE problem timespan is modified to be sampled uniformly at random from the interval (t2 - end_time_regularization, t2 + end_time_regularization).

```
In [10]: abstract type NeuralODELayer <: LuxCore.AbstractExplicitContainerLayer{(:model,)} end
basic_tgrad(u,p,t) = zero(u)

struct NeuralODETimeRegularized{M,P,RE,T,ETR,A,K} <: NeuralODELayer
model::M
p::P
re::RE
tspan::T
end_time_regularization::ETR
args::A
kwargs::K
end

function NeuralODETimeRegularized(model,tspan,end_time_regularization,args...;p = nothing,kwargs...)
_p,re = Flux.destructure(model)
if p == nothing
p = _p
end
NeuralODETimeRegularized{typeof(model),typeof(p),typeof(re),
typeof(tspan), typeof(end_time_regularization), typeof(args),typeof(kwargs)}(
model,p,re,tspan,end_time_regularization,args,kwargs)
end

function (n::NeuralODETimeRegularized)(x,p,n.p)
dudt_(u,p,t) = n.re(p)(u)
ff = ODEFunction{false}(dudt_,tgrad=basic_tgrad)

# uniformly at random sample the endpoint of the ODE problem timespan from
# the interval (t2 - end_time_regularization, t2 + end_time_regularization).
perturbation = convert{typeof(n.tspan[1])}(rand{Uniform{(-n.end_time_regularization, n.end_time_regularization)}})
prob = ODEProblem{false}(ff,x,(n.tspan[1], n.tspan[2]+perturbation),p)

# solve ODE
sense = InterpolatingAdjoint(autojacvec=ZygoteVJP())
sol = solve(prob,n.args...;sensealg=sense,n.kwargs...)
return sol
end
```

Neural ODE variant: Second Order NODE

Similar to the above section, but we implement a second-order NODE. The main modification is that the input is treated as a concatenation of the initial conditions for du_0 and u_0 , and we solve a second-order ODE $u'' = NN(u)$. The output contains final trajectory values for both u and du .

```
In [11]: struct NeuralODESecondOrder{M,P,RE,T,A,K} <: NeuralODELayer
  model::M
  p::P
  re::RE
  tspan::T
  args::A
  kwargs::K
end

function NeuralODESecondOrder(model,tspan,args...;p = nothing,kwargs...)
  _p,re = Flux.destructure(model)
  if p == nothing
    p = _p
  end
  NeuralODESecondOrder(typeof(model),typeof(p),typeof(re),
    typeof(tspan),typeof(args),typeof(kwargs)) {
    model,p,re,tspan,args,kwargs}
end

function (n::NeuralODESecondOrder)(x,p=n.p)
  # second order ODE definition
  ddudt_(du,u,p,t) = n.re(p)(u)

  # split input into initial conditions for du0 and u0
  nn_dimension = convert{Int64, size(x)[1]/2}
  du0 = x[1:nn_dimension, :]
  u0 = x[nn_dimension+1:end, :]

  # solve second order ODE using given initial conditions
  prob = SecondOrderODEProblem{false}(ddudt_,du0,u0,getfield(n,:tspan),p)
  sol = solve(prob,n.args...;p=n.p,kwargs...)

  # output is a concatenation of the final trajectory values for u and du
  return Array(sol)
end
```

ANODE training function for MNIST classification

This is the main function for training an ANODE for MNIST classification, with all possible variants (zero-padding augmentation, input layer augmentation, end time regularization, and second order NODEs). Note that the first layer reshapes the image into a vector of length 784, and the final output layer is a fully connected layer with 10 output nodes (corresponding to the one-hot encoding of the 10 digit classes). The neural ODE structure varies based on the optional inputs for the different variants.

```
In [253]: function train_mnist_anode(augment_dims; input_layer_augmentation=false, end_time_regularization=0, second_order=false, get_test_acc=true)
  # Given number of augmentation dimensions, train ANODE. (If augment_dims=0, this is just regular NODE.)
  # Optional arguments are for ANODE variants: input layer augmentation, end time regularization, and second order ANODEs.
  # Outputs: classification accuracies and losses over training iterations.

  # reshape image into Length-784 vector
  reshape_layer = x->reshape(x,(784,:))

  # input layer maps image to lower dimensional space, and we add augmenting dimensions
  if input_layer_augmentation
    input_layer = Flux.Dense(784,20+augment_dims,tanh) # train network to learn the augmented dimensions
  elseif second_order
    # pad the two halves of the downsampled input (corresponding to u and du)
    input_layer = Flux.Chain(Flux.Dense(784,80,tanh),
      x->cat(dims=1, x[1:40, :], zeros(Float32, augment_dims, bs), x[41:80, :], zeros(Float32, augment_dims, bs)))
  else
    input_layer = Flux.Chain(Flux.Dense(784,20,tanh),
      x->cat(dims=1, x, zeros(Float32, augment_dims, bs))) # pad with augment_dims zeros
  end

  # different neural network architectures for first order vs second order ODEs
  if second_order
    nn = Flux.Chain(
      Flux.Dense(40+augment_dims, 20,tanh),
      Flux.Dense(20, 20, tanh),
      Flux.Dense(20, 40+augment_dims,tanh),
    )
    # final output is 10-dimensional since there are 10 classes of digits
    output_layer = Flux.Chain(x->x[1:40, :], Flux.Dense(40,10))
  else
    nn = Flux.Chain(
      Flux.Dense(20+augment_dims, 10,tanh),
      Flux.Dense(10, 10,tanh),
      Flux.Dense(10, 20+augment_dims,tanh),
    )
    output_layer = Flux.Chain(x->x[1:20, :], Flux.Dense(20,10))
  end

  # different Neural ODE variants defined earlier
  if end_time_regularization != 0
    nn_ode = NeuralODETimeRegularized(nn, (0.f0, 1.f0), end_time_regularization, Tsit5(), save_everystep = false, reitol = 1e-3, abstol = 1e-3, save_start = false)
  elseif second_order
    nn_ode = NeuralODESecondOrder(nn, (0.f0, 1.f0), Tsit5(), save_everystep = false, reitol = 1e-3, abstol = 1e-3, save_start = false)
  else
    nn_ode = NeuralODE(nn, (0.f0, 1.f0), Tsit5(), save_everystep = false, reitol = 1e-3, abstol = 1e-3, save_start = false)
  end

  # create final model by chaining together layers
  if second_order
    m = Flux.Chain(reshape_layer, input_layer, nn_ode, output_layer)
  else
    m = Flux.Chain(reshape_layer, input_layer, nn_ode, DiffEqArrayTo_Array, output_layer)
  end

  loss(x,y) = logitcrossentropy(m(x),y)

  accuracies = []
  losses = []

  cb_mnist() = begin
    # compute accuracies and losses over test dataset or train dataset
    if get_test_acc
      append!(accuracies, accuracy(m, zip(x_test,y_test)))
      append!(losses, total_loss(m, zip(x_test,y_test)))
    else
      append!(accuracies, accuracy(m, zip(x_train,y_train)))
      append!(losses, total_loss(m, zip(x_train,y_train)))
    end
  end

  for i in 1:5 # training epochs
    Flux.train!(loss, Flux.params(m), zip(x_train, y_train), opt, cb = cb_mnist)
  end

  print("Final accuracy: ", accuracies[end])
  return accuracies, losses
end
```

Out[253]: train_mnist_anode (generic function with 1 method)

Train regular Neural ODE

```
In [102]: node_accuracies, node_losses = train_mnist_anode(0)
```

Final accuracy: 0.732

```
Out[102]: (Any[0.227, 0.399, 0.443, 0.463, 0.49, 0.538, 0.588, 0.601, 0.586, 0.633 ..., 0.728, 0.726, 0.722, 0.735, 0.739, 0.745, 0.749, 0.738, 0.722, 0.732], Any[26.773666f0, 21.134886f0, 19.157936f0, 16.258987f0, 15.254635f0, 14.603728f0, 13.128592f0, 12.664176f0, 12.627069f0, 11.540196f0 ..., 8.321392f0, 8.458588f0, 8.439966f0, 8.276887f0, 8.009961f0, 7.924491f0, 7.8128023f0, 8.056908f0, 8.4869175f0, 8.207112f0])
```

Zero-padding augmentation

```
In [122]: anode_accuracies_10, anode_losses_10 = train_mnist_anode(10)
```

Final accuracy: 0.77

```
Out[122]: (Any[0.235, 0.398, 0.483, 0.499, 0.519, 0.534, 0.566, 0.613, 0.634, 0.617 ..., 0.759, 0.77, 0.782, 0.796, 0.786, 0.776, 0.786, 0.783, 0.781, 0.77], Any[25.475597f0, 20.12066f0, 16.278515f0, 14.906236f0, 14.202824f0, 13.492475f0, 12.7209835f0, 11.868252f0, 11.405236f0, 11.339441f0 ..., 7.768994f0, 7.4147224f0, 7.111907f0, 6.965706f0, 7.059421f0, 7.1792564f0, 7.217086f0, 7.229756f0, 7.2592144f0, 7.7979345f0])
```

```
In [124]: anode_accuracies_25, anode_losses_25 = train_mnist_anode(25)
```

Final accuracy: 0.798

```
Out[124]: (Any[0.279, 0.233, 0.493, 0.569, 0.605, 0.624, 0.627, 0.666, 0.703, 0.708 ..., 0.764, 0.756, 0.754, 0.77, 0.795, 0.8, 0.803, 0.802, 0.804, 0.798], Any[23.411448f0, 20.532196f0, 15.624832f0, 13.48805f0, 12.494654f0, 11.721497f0, 11.159152f0, 10.257824f0, 9.38868f0, 8.910438f0 ..., 8.114785f0, 8.329447f0, 8.026104f0, 7.455347f0, 6.857174f0, 6.432046f0, 6.2200937f0, 6.039778f0, 5.956437f0, 6.135666f0])
```


```
In [128]: anode_accuracies_40, anode_losses_40 = train_mnist_anode(40)
```

Final accuracy: 0.811

```
Out[128]: (Any[0.297, 0.422, 0.585, 0.592, 0.586, 0.633, 0.651, 0.652, 0.686, 0.69 ..., 0.766, 0.769, 0.778, 0.785, 0.797, 0.801, 0.805, 0.815, 0.821, 0.811], Any[21.961197f0, 16.545862f0, 14.155242f0, 12.974012f0, 12.207743f0, 10.956065f0, 10.84268f0, 10.399176f0, 9.495364f0, 9.2334385f0 ..., 6.64111f0, 6.661842f0, 6.435278f0, 6.2573957f0, 6.122465f0, 5.992382f0, 5.9030104f0, 5.742066f0, 5.6628275f0, 5.8586f0])
```


```
In [188]: plot(1:50, node_accuracies, label="NODE")
plot!(1:50, anode_accuracies_10, label="ANODE (aug dim = 10)")
plot!(1:50, anode_accuracies_25, label="ANODE (aug dim = 25)")
plot!(1:50, anode_accuracies_40, label="ANODE (aug dim = 40)")
title("ANODE Test Accuracy on MNIST classification")
xlabel!("Iterations")
ylabel!("Test Accuracy")
```

Out[188]: ANODE Test Accuracy on MNIST classification



```
In [189]: plot(1:50, node_losses, label="NODE")
plot!(1:50, anode_losses_10, label="ANODE (aug dim = 10)")
plot!(1:50, anode_losses_25, label="ANODE (aug dim = 25)")
plot!(1:50, anode_losses_40, label="ANODE (aug dim = 40)")
title("ANODE Test Loss on MNIST classification")
xlabel!("Iterations")
ylabel!("Cross-Entropy Loss")
```

Out[189]: ANODE Test Loss on MNIST classification



Input-layer augmentation

```
In [138]: il_anode_accuracies_10, il_anode_losses_10 = train_mnist_anode(10, input_layer_augmentation=true)
```

Final accuracy: 0.807

```
Out[138]: (Any[0.236, 0.342, 0.461, 0.477, 0.556, 0.618, 0.614, 0.63, 0.615, 0.61 ..., 0.774, 0.798, 0.815, 0.813, 0.806, 0.795, 0.804, 0.813, 0.815, 0.807], Any[26.976202f0, 19.111366f0, 15.211487f0, 13.73826f0, 11.919678f0, 11.102933f0, 10.682887f0, 10.745924f0, 10.585981f0, 10.429465f0 ..., 7.3744354f0, 6.7953806f0, 6.325192f0, 6.3474174f0, 6.4996414f0, 6.5120573f0, 6.21073f0, 5.9839854f0, 5.9134407f0, 6.1013794f0])
```

```
In [132]: il_anode_accuracies_25, il_anode_losses_25 = train_mnist_anode(25, input_layer_augmentation=true)
```

Final accuracy: 0.814

```
Out[132]: (Any[0.339, 0.349, 0.481, 0.502, 0.555, 0.58, 0.603, 0.683, 0.681, 0.654 ..., 0.804, 0.797, 0.797, 0.791, 0.798, 0.803, 0.813, 0.814, 0.81, 0.814], Any[24.32352f0, 18.394129f0, 14.81003f0, 13.696833f0, 12.909499f0, 12.173111f0, 11.961259f0, 10.760969f0, 10.16151f0, 9.747705f0 ..., 6.1982403f0, 6.366974f0, 6.520851f0, 6.6333575f0, 6.5653963f0, 6.317669f0, 6.182312f0, 6.120674f0, 6.1274567f0, 6.1962147f0])
```


```
In [136]: il_anode_accuracies_40, il_anode_losses_40 = train_mnist_anode(40, input_layer_augmentation=true)
```

Final accuracy: 0.823

```
Out[136]: (Any[0.246, 0.404, 0.495, 0.57, 0.594, 0.592, 0.614, 0.638, 0.678, 0.672 ..., 0.799, 0.803, 0.806, 0.816, 0.809, 0.809, 0.819, 0.83, 0.825, 0.823], Any[21.776491f0, 18.317554f0, 15.384143f0, 13.855808f0, 12.930713f0, 12.045832f0, 11.108894f0, 10.464067f0, 10.092501f0, 10.068508f0 ..., 6.5776024f0, 6.4511743f0, 6.0879784f0, 6.0418367f0, 6.0769715f0, 6.0826015f0, 6.0280097f0, 5.6082563f0, 5.6620355f0, 5.829573f0])
```


```
In [190]: plot(1:50, node_accuracies, label="NODE")
plot(1:50, anode_accuracies_10, label="ANODE (aug dim = 10)")
plot(1:50, il_anode_accuracies_10, label="IL-ANODE (aug dim = 10)")
plot(1:50, il_anode_accuracies_25, label="IL-ANODE (aug dim = 25)")
plot(1:50, il_anode_accuracies_40, label="IL-ANODE (aug dim = 40)")
title("IL-ANODE Test Accuracy on MNIST classification")
xlabel("Iterations")
ylabel("Test Accuracy")
```

Out[190]: IL-ANODE Test Accuracy on MNIST classification



```
In [191]: plot(1:50, node_losses, label="NODE")
plot(1:50, anode_losses_10, label="ANODE (aug dim = 10)")
plot(1:50, il_anode_losses_10, label="IL-ANODE (aug dim = 10)")
plot(1:50, anode_losses_25, label="IL-ANODE (aug dim = 25)")
plot(1:50, anode_losses_40, label="IL-ANODE (aug dim = 40)")
title("IL-ANODE Test Loss on MNIST classification")
xlabel("Iterations")
ylabel("Cross-Entropy Loss")
```

Out[191]: IL-ANODE Test Loss on MNIST classification



Temporal regularization

```
In [165]: reg_anode_accuracies_20, reg_anode_losses_20 = train_mnist_anode(10, end_time_regularization=0.2)
```

Final accuracy: 0.81

```
Out[165]: (Any[0.424, 0.512, 0.571, 0.626, 0.633, 0.645, 0.657, 0.667, 0.692, 0.705 ..., 0.812, 0.814, 0.812, 0.809, 0.809, 0.802, 0.795, 0.81, 0.806, 0.81], Any[19.747875f0, 17.554031f0, 14.786218f0, 12.163734f0, 11.380127f0, 1
0.914139f0, 10.48744f0, 10.126495f0, 9.667911f0, 9.291637f0 ..., 5.919667f0, 6.081716f0, 6.2100263f0, 6.3855615f0, 6.3893394f0, 6.420374f0, 6.322182f0, 6.232897f0, 6.1347775f0, 6.098957f0])
```

```
In [166]: reg_anode_accuracies_30, reg_anode_losses_30 = train_mnist_anode(10, end_time_regularization=0.3)
```

Final accuracy: 0.808

```
Out[166]: (Any[0.357, 0.362, 0.533, 0.575, 0.602, 0.611, 0.63, 0.656, 0.686, 0.697 ..., 0.79, 0.794, 0.797, 0.805, 0.813, 0.814, 0.809, 0.809, 0.81, 0.808], Any[21.559385f0, 19.332924f0, 15.898358f0, 13.9182625f0, 12.294415f0, 1
1.268187f0, 10.86916f0, 10.087314f0, 9.530009f0, 9.173318f0 ..., 6.363705f0, 6.268433f0, 6.068194f0, 5.8814287f0, 5.7526555f0, 5.735334f0, 5.718261f0, 5.8403835f0, 5.9131985f0, 5.9233027f0])
```


```
In [181]: reg_anode_accuracies_60, reg_anode_losses_60 = train_mnist_anode(10, end_time_regularization=0.6)
```

Final accuracy: 0.821

```
Out[181]: (Any[0.246, 0.522, 0.58, 0.556, 0.565, 0.592, 0.643, 0.702, 0.697, 0.693 ..., 0.777, 0.781, 0.79, 0.791, 0.798, 0.8, 0.804, 0.825, 0.821, 0.821], Any[24.08501f0, 15.513119f0, 13.363345f0, 12.5917845f0, 12.221241f0, 11.4
7541f0, 10.416765f0, 9.56783f0, 9.439296f0, 9.36563f0 ..., 7.383833f0, 7.5351725f0, 7.1987896f0, 7.0500393f0, 6.9000525f0, 7.007737f0, 6.736189f0, 6.458054f0, 6.536541f0, 6.677477f0])
```


```
In [195]: plot(1:50, node_accuracies, label="NODE")
plot(1:50, anode_accuracies_10, label="ANODE (aug dim = 10)")
plot(1:50, reg_anode_accuracies_20, label="time-regularized ANODE (aug dim = 10, b = 0.2)")
plot(1:50, reg_anode_accuracies_30, label="time-regularized ANODE (aug dim = 10, b = 0.3)")
plot(1:50, reg_anode_accuracies_60, label="time-regularized ANODE (aug dim = 10, b = 0.6)")
title("Time-regularized ANODE Test Accuracy on MNIST")
xlabel("Iterations")
ylabel("Test Accuracy")
```

Out[195]: Time-regularized ANODE Test Accuracy on MNIST



```
In [217]: plot(1:50, node_losses, label="NODE")
plot(1:50, anode_losses_10, label="ANODE (aug dim = 10)")
plot(1:50, reg_anode_losses_20, label="time-regularized ANODE (aug dim = 10, b = 0.2)")
plot(1:50, reg_anode_losses_30, label="time-regularized ANODE (aug dim = 10, b = 0.3)")
plot(1:50, reg_anode_losses_60, label="time-regularized ANODE (aug dim = 10, b = 0.6)")
title("Time-regularized ANODE Test Loss on MNIST")
xlabel("Iterations")
ylabel("Cross-Entropy Loss")
```

Out[217]: Time-regularized ANODE Test Loss on MNIST



We also compare the train accuracies and losses to see the effect of regularization.

```
In [203]: train_anode_accuaries_10, train_anode_losses_10 = train_mnist_anode(10, get_test_acc=false)
```

Final accuracy: 0.946

```
Out[203]: (Any[0.241, 0.38, 0.498, 0.604, 0.67, 0.678, 0.651, 0.699, 0.721, 0.754 ... 0.918, 0.914, 0.917, 0.937, 0.94, 0.934, 0.935, 0.934, 0.942, 0.946], Any[23.077784f0, 20.176426f0, 14.558219f0, 12.5663185f0, 10.771791f0, 9.7211075f0, 9.553288f0, 8.889776f0, 8.6946335f0, 7.9237175f0 ... 2.532746f0, 2.6941848f0, 2.579479f0, 2.269286f0, 2.2287428f0, 2.3123062f0, 2.4240556f0, 2.3259807f0, 2.1167514f0, 2.0490487f0])
```

```
In [212]: train_reg_anode_accuaries_20, train_reg_anode_losses_20 = train_mnist_anode(10, end_time_regularization=0.2, get_test_acc=false)
```

Final accuracy: 0.938

```
Out[212]: (Any[0.111, 0.41, 0.462, 0.5, 0.498, 0.521, 0.581, 0.612, 0.618, 0.653 ... 0.919, 0.915, 0.914, 0.919, 0.925, 0.925, 0.933, 0.937, 0.938], Any[30.930473f0, 17.823364f0, 15.5176525f0, 15.553515f0, 14.845526f0, 13.960464f0, 12.671738f0, 11.6204605f0, 10.699567f0, 9.843736f0 ... 3.1395802f0, 3.1099803f0, 3.0403104f0, 2.9484227f0, 2.7943244f0, 2.6481507f0, 2.5220902f0, 2.4438884f0, 2.3126972f0, 2.3128533f0])
```

```
In [199]: train_reg_anode_accuaries_30, train_reg_anode_losses_30 = train_mnist_anode(10, end_time_regularization=0.3, get_test_acc=false)
```

Final accuracy: 0.904

```
Out[199]: (Any[0.197, 0.304, 0.386, 0.484, 0.549, 0.597, 0.592, 0.583, 0.581, 0.597 ... 0.885, 0.883, 0.88, 0.885, 0.893, 0.902, 0.908, 0.915, 0.923, 0.904], Any[31.288296f0, 22.87244f0, 18.19588f0, 15.995473f0, 14.732131f0, 13.929122f0, 13.421138f0, 12.851914f0, 12.369618f0, 11.980809f0 ... 3.9037642f0, 3.8571222f0, 3.7724004f0, 3.7433622f0, 3.7173815f0, 3.5264835f0, 3.349937f0, 3.0599642f0, 2.924942f0, 3.1416585f0])
```


```
In [207]: train_reg_anode_accuaries_60, train_reg_anode_losses_60 = train_mnist_anode(10, end_time_regularization=0.6, get_test_acc=false)
```

Final accuracy: 0.926

```
Out[207]: (Any[0.182, 0.455, 0.51, 0.557, 0.636, 0.623, 0.619, 0.67, 0.706, 0.719 ... 0.906, 0.901, 0.907, 0.906, 0.916, 0.922, 0.928, 0.932, 0.929, 0.926], Any[25.784758f0, 18.008217f0, 14.987472f0, 12.914676f0, 11.863629f0, 11.472921f0, 10.8568535f0, 10.073899f0, 9.285956f0, 8.785873f0 ... 3.2794962f0, 3.2985308f0, 3.3525686f0, 3.2584991f0, 3.0158727f0, 2.7157683f0, 2.5784726f0, 2.5748463f0, 2.6814878f0, 2.6752093f0])
```


```
In [215]: plot(1:50, train_anode_accuaries_10, label="ANODE (aug dim = 10)", ylims=(0, 1))
plot(1:50, train_reg_anode_accuaries_20, label="time-regularized ANODE (aug dim = 10, b = 0.2)")
plot(1:50, train_reg_anode_accuaries_30, label="time-regularized ANODE (aug dim = 10, b = 0.3)")
plot(1:50, train_reg_anode_accuaries_60, label="time-regularized ANODE (aug dim = 10, b = 0.6)")
title("Time-regularized ANODE Train Accuracy on MNIST")
xlabel("Iterations")
ylabel("Train Accuracy")
```

Out[215]: Time-regularized ANODE Train Accuracy on MNIST



```
In [219]: plot(1:50, train_anode_losses_10, label="ANODE (aug dim = 10)")
plot(1:50, train_reg_anode_losses_20, label="time-regularized ANODE (aug dim = 10, b = 0.2)")
plot(1:50, train_reg_anode_losses_30, label="time-regularized ANODE (aug dim = 10, b = 0.3)")
plot(1:50, train_reg_anode_losses_60, label="time-regularized ANODE (aug dim = 10, b = 0.6)")
title("Time-regularized ANODE Train Loss on MNIST")
xlabel("Iterations")
ylabel("Cross-Entropy Loss")
```

Out[219]: Time-regularized ANODE Train Loss on MNIST




Second order NODEs

```
In [274]: second_order_node_accuracies_0, second_order_node_losses_0 = train_mnist_anode(0, second_order=true)

Final accuracy: 0.554

Out[274]: (Any[0.109, 0.088, 0.116, 0.122, 0.226, 0.291, 0.309, 0.334, 0.379, 0.372 ... 0.529, 0.526, 0.541, 0.537, 0.553, 0.538, 0.528, 0.532, 0.549, 0.554], Any[41.822056f0, 39.741383f0, 27.164011f0, 24.48871f0, 25.409086f0, 2
4.552261f0, 22.998236f0, 21.447826f0, 19.802427f0, 18.352165f0 ... 14.319466f0, 14.668753f0, 14.380504f0, 13.760771f0, 13.430332f0, 13.577651f0, 13.909873f0, 13.633666f0, 13.377003f0, 13.244428f0])

In [284]: plot(1:50, node_accuracies, label="NODE")
plot(1:50, anode_accuracies_10, label="ANODE (aug dim = 10)")
plot(1:50, second_order_node_accuracies_0, label="second-order NODE")
title("Second-order ANODE Test Accuracy on MNIST")
xlabel("Iterations")
ylabel("Test Accuracy")
```



```
In [285]: plot(1:50, node_losses, label="NODE")
plot(1:50, anode_losses_10, label="ANODE (aug dim = 10)")
plot(1:50, second_order_node_losses_0, label="second-order NODE")
title("Second-order ANODE Test Loss on MNIST")
xlabel("Iterations")
ylabel("Cross-Entropy Loss")
```

